
Principles of Distributed Database Systems

TS. Phan Thị Hà

Outline

- Introduction
- Distributed and Parallel Database Design
- Distributed Data Control
- Distributed Query Processing
- Distributed Transaction Processing
- Data Replication
- Database Integration – Multidatabase Systems
- Parallel Database Systems
- Peer-to-Peer Data Management
- **Big Data Processing**
- NoSQL, NewSQL and Polystores
- Web Data Management

Outline

- Big Data Processing
 - ❑ Distributed storage systems
 - ❑ Processing platforms
 - ❑ Stream data management
 - ❑ Graph analytics
 - ❑ Data lake

Four Vs

■ Volume

- ▣ Increasing data size: petabytes (10^{15}) to zettabytes (10^{21})

■ Variety

- ▣ Multimodal data: structured, images, text, audio, video
- ▣ 90% of currently generated data unstructured

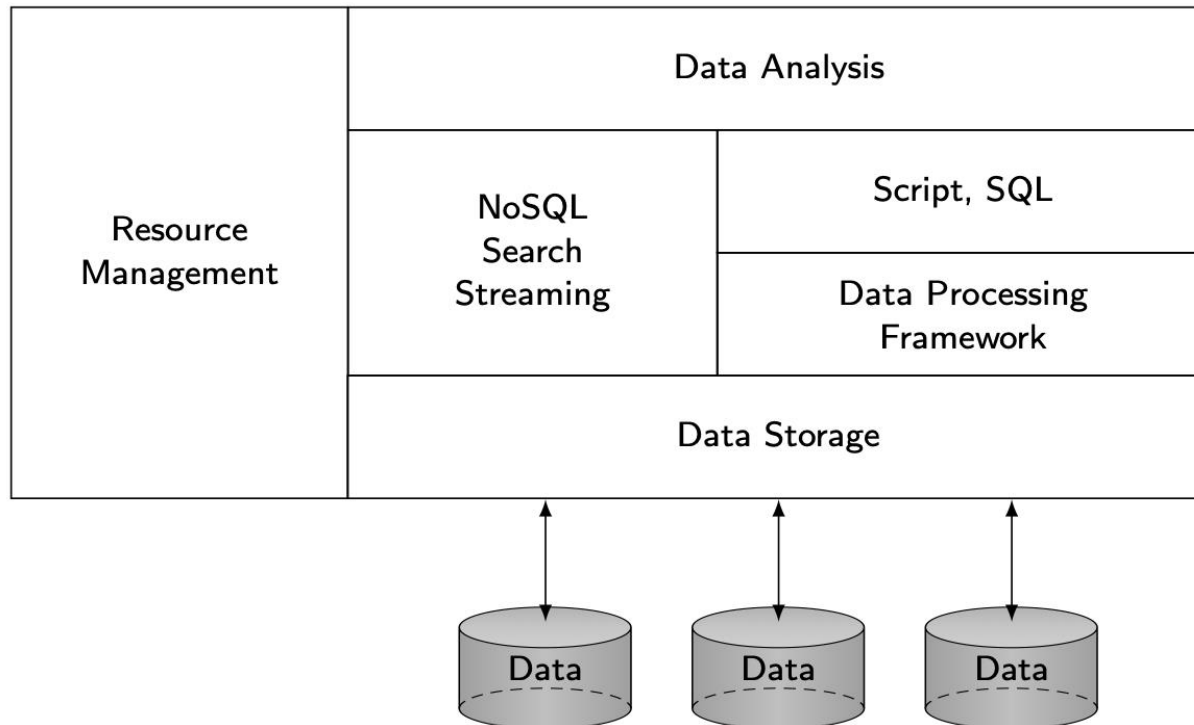
■ Velocity

- ▣ Streaming data at high speed
- ▣ Real-time processing

■ Veracity

- ▣ Data quality

Big Data Software Stack



Outline

- Big Data Processing
 - Distributed storage systems
 - Processing platforms
 - Stream data management
 - Graph analytics

Distributed Storage System

Storing and managing data across the nodes of a shared-nothing cluster

■ Object-based

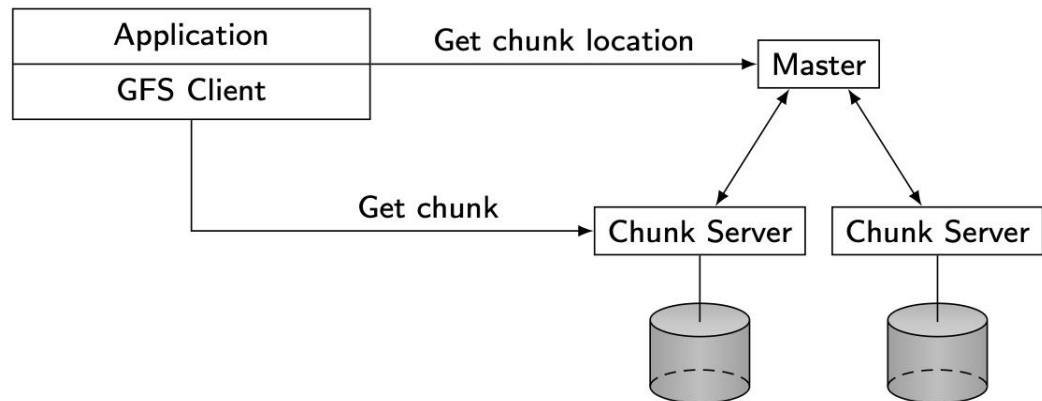
- ❑ Object = $\langle \text{oid}, \text{data}, \text{metadata} \rangle$
- ❑ Metadata can be different for different object
- ❑ Easy to move
- ❑ Flat object space \rightarrow billions/trillions of objects
- ❑ Easily accessed through REST-based API (get/put)
- ❑ Good for high number of small objects (photos, mail attachments)

■ File-based

- ❑ Data in files of fixed- or variable-length records
- ❑ Metadata-per-file stored separately from file
- ❑ For large data, a file needs to be partitioned and distributed

Google File System (GFS)

- Targets shared-nothing clusters of thousands of machines
- Targets applications with characteristics:
 - ❑ Very large files (several gigabytes)
 - ❑ Mostly read and append workloads
 - ❑ High throughput more important than low latency
- Interface: create, open, read, write, close, delete, snapshot, record append



Outline

- Big Data Processing
 - ❑ Distributed storage systems
 - ❑ Processing platforms
 - ❑ Stream data management
 - ❑ Graph analytics

Big Data Processing Platforms

- Applications that do not need full DBMS functionality
 - Data analysis of very large data sets
 - Highly dynamic, irregular, schemaless, ...
- “Embarrassingly parallel problems”
- MapReduce/Spark
- Advantages
 - Flexibility
 - Scalability
 - Efficiency
 - Fault-tolerance
- Disadvantage
 - Reduced functionality
 - Increased programming effort

MapReduce Basics

- Simple programming model
 - Data structured as (key, value) pairs
 - E.g. (doc-id, content); (word, count)
 - Functional programming style with two functions
 - $\text{map}(k1, v1) \rightarrow \text{list}(k2, v2)$
 - $\text{reduce}(k2, \text{list}(v2)) \rightarrow \text{list}(v3)$
- Implemented on a distributed file system (e.g. Google File System) on very large clusters

map Function

- User-defined function
 - ❑ Processes input (key, value) pairs
 - ❑ Produces a set of **intermediate** (key, value) pairs
 - ❑ Executes on multiple machines (called **mapper**)
- map function I/O
 - ❑ **Input:** read a **chunk** from distributed file system (DFS)
 - ❑ **Output:** Write to intermediate file on local disk
- MapReduce library
 - ❑ Execute map function
 - ❑ Groups together all intermediate values with same key
 - ❑ Passes these lists to reduce function
- Effect of map function
 - ❑ Processes and partitions input data
 - ❑ Builds a distributed map (transparent to user)
 - ❑ Similar to “group by” operator in SQL

reduce Function

- User-defined function
 - ❑ Accepts one intermediate key and a set of values for that key (i.e. a list)
 - ❑ Merges these values together to form a (possibly) smaller set
 - ❑ Computes the reduce function generating, typically, zero or one output per invocation
 - ❑ Executes on multiple machines (called **reducer**)
- reduce function I/O
 - ❑ **Input:** read from intermediate files using remote reads on local files of corresponding mappers
 - ❑ **Output:** Write result back to DFS
- Effect of map function
 - ❑ Similar to aggregation function in SQL

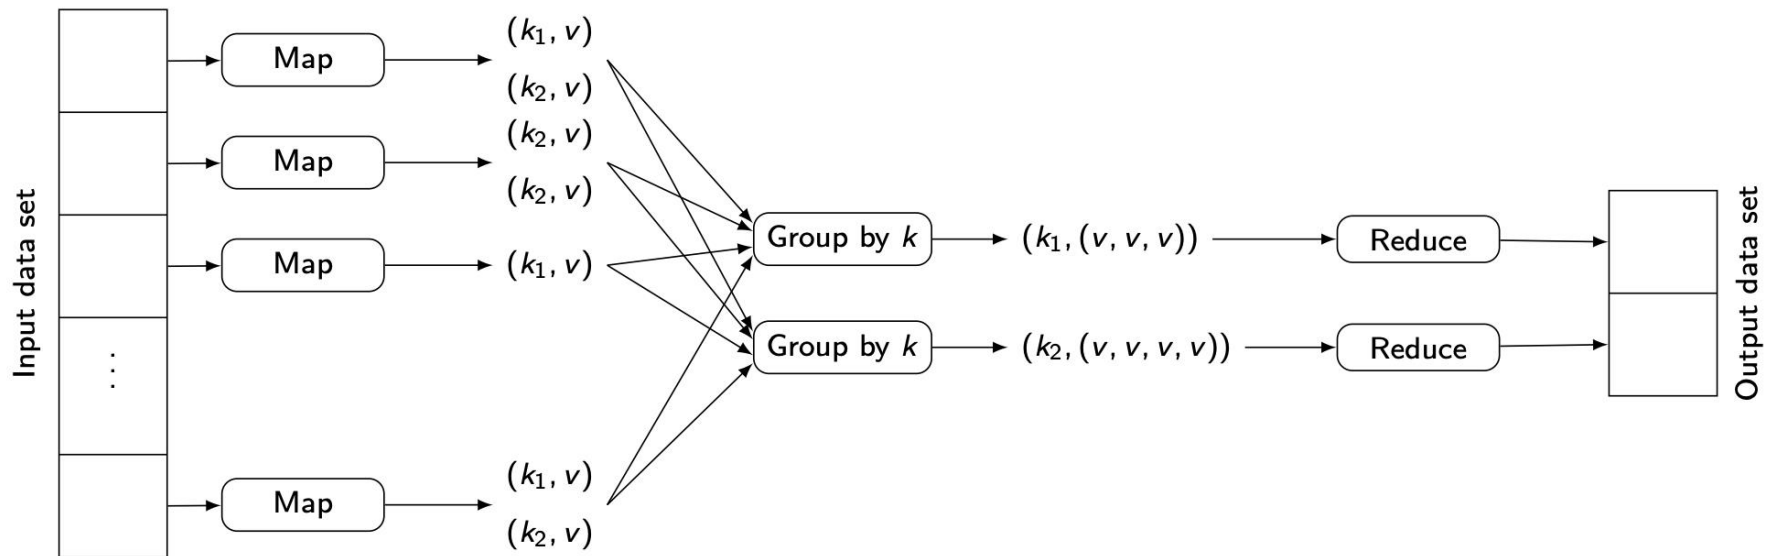
Example

Consider EMP (ENO, ENAME, TITLE, CITY)

```
SELECT      CITY, COUNT (*)  
FROM        EMP  
WHERE        ENAME LIKE "%Smith"  
GROUP BY    CITY
```

```
map (Input: (TID,EMP), Output: (CITY, 1)  
    if EMP.ENAME like ``\%Smith'' return (CITY, 1)  
reduce (Input: (CITY, list(1)), Output: (CITY,  
SUM(list)))  
    return (CITY, SUM(1))
```

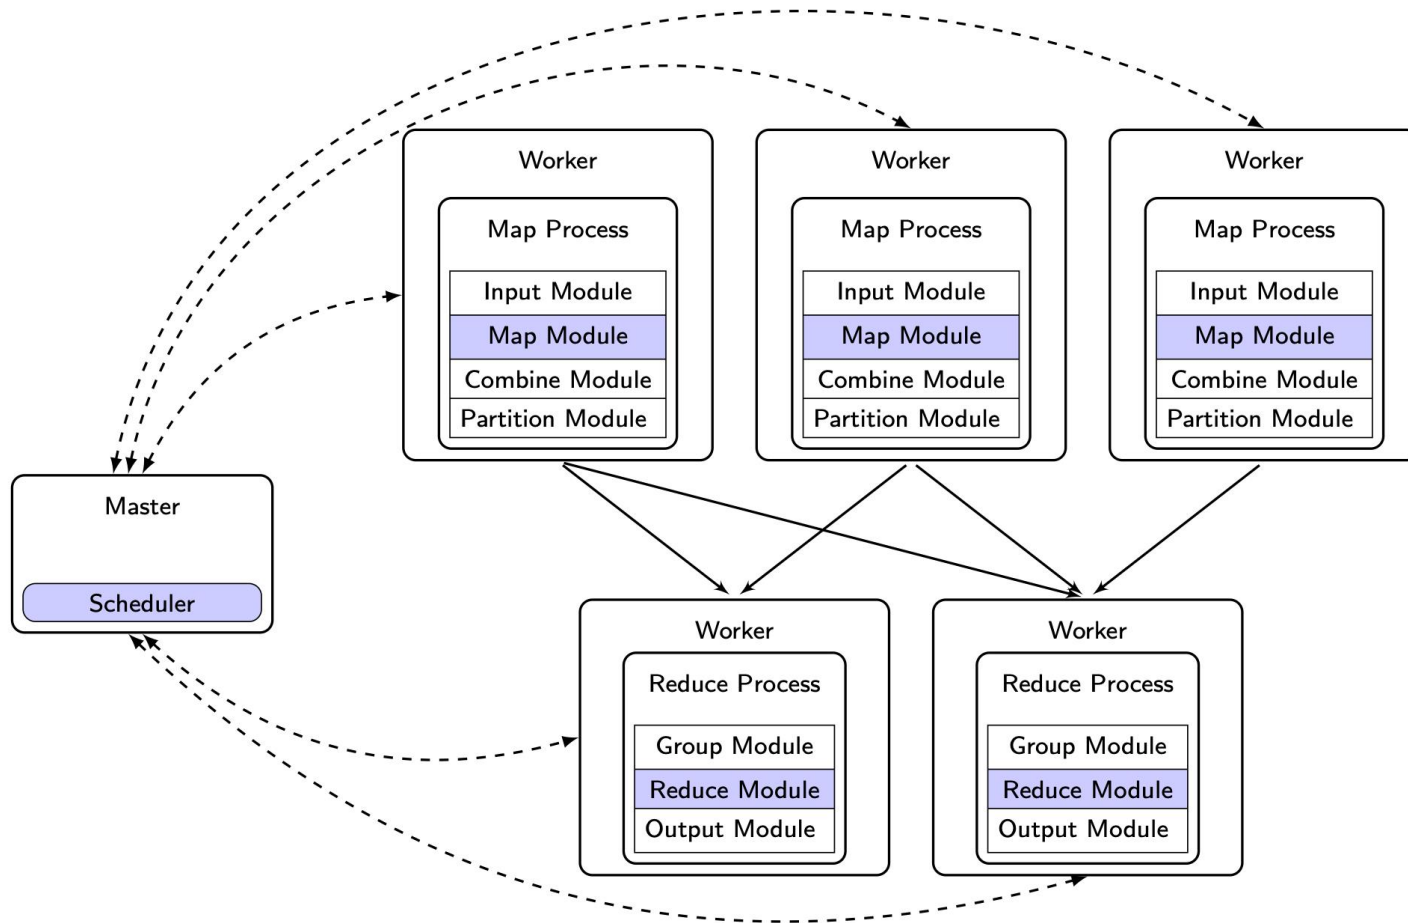
MapReduce Processing



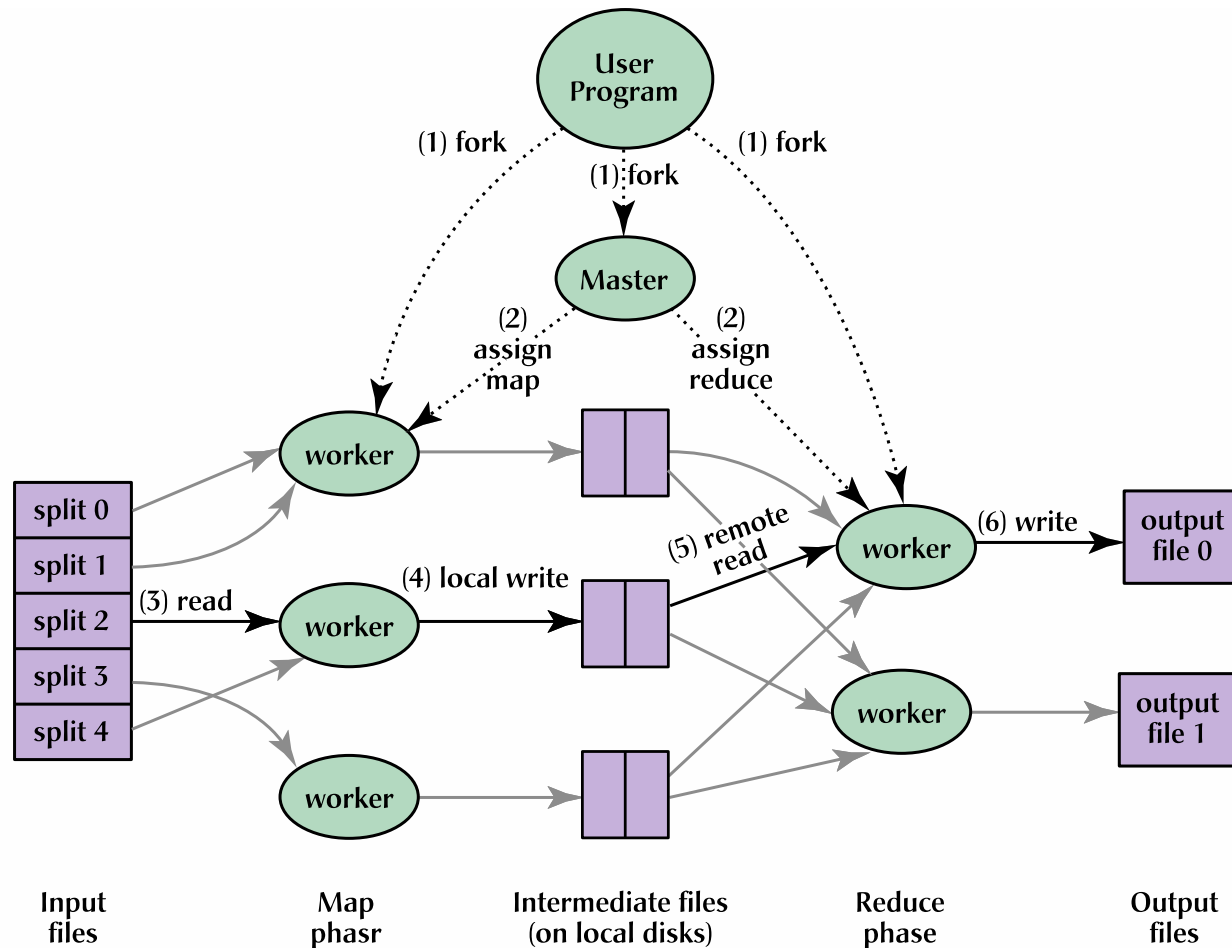
Hadoop Stack

Yarn	Third party analysis tools R (statistics), Mahout (machine learning), ...	
	Hbase	Hive & HiveQL
		Hadoop (MapReduce engine)
	Hadoop Distributed File System (HDFS)	

Master-Worker Architecture

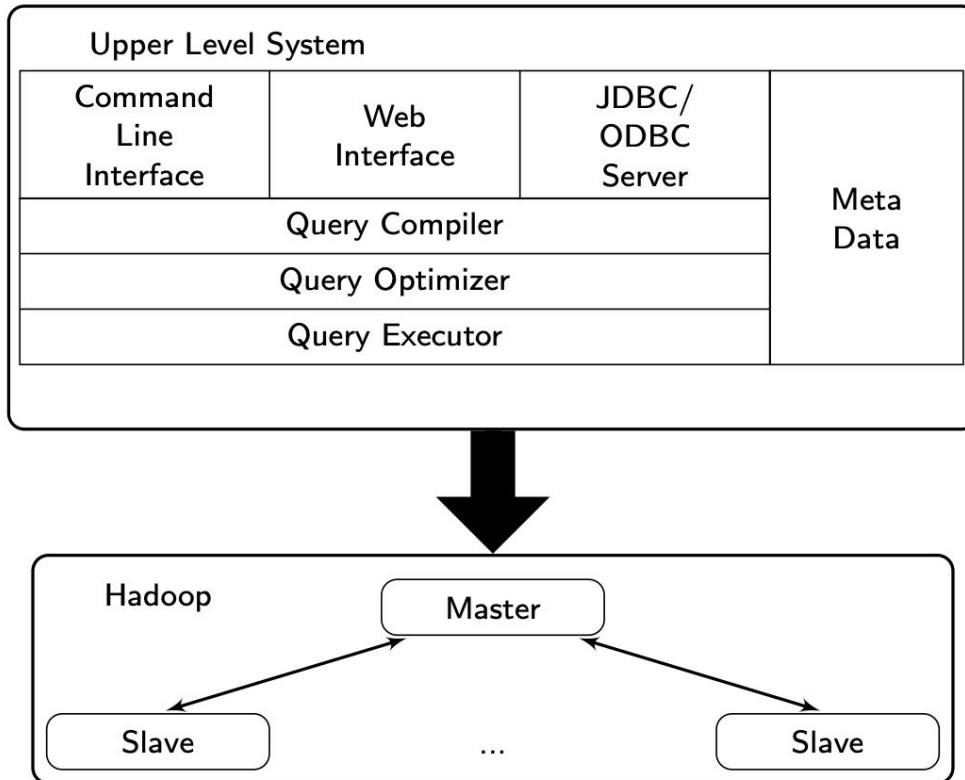


Execution Flow



From: J. Dean and S. Ghemawat. MapReduce: Simplified data processing on large clusters, *Comm. ACM*, 51(1), 2008.

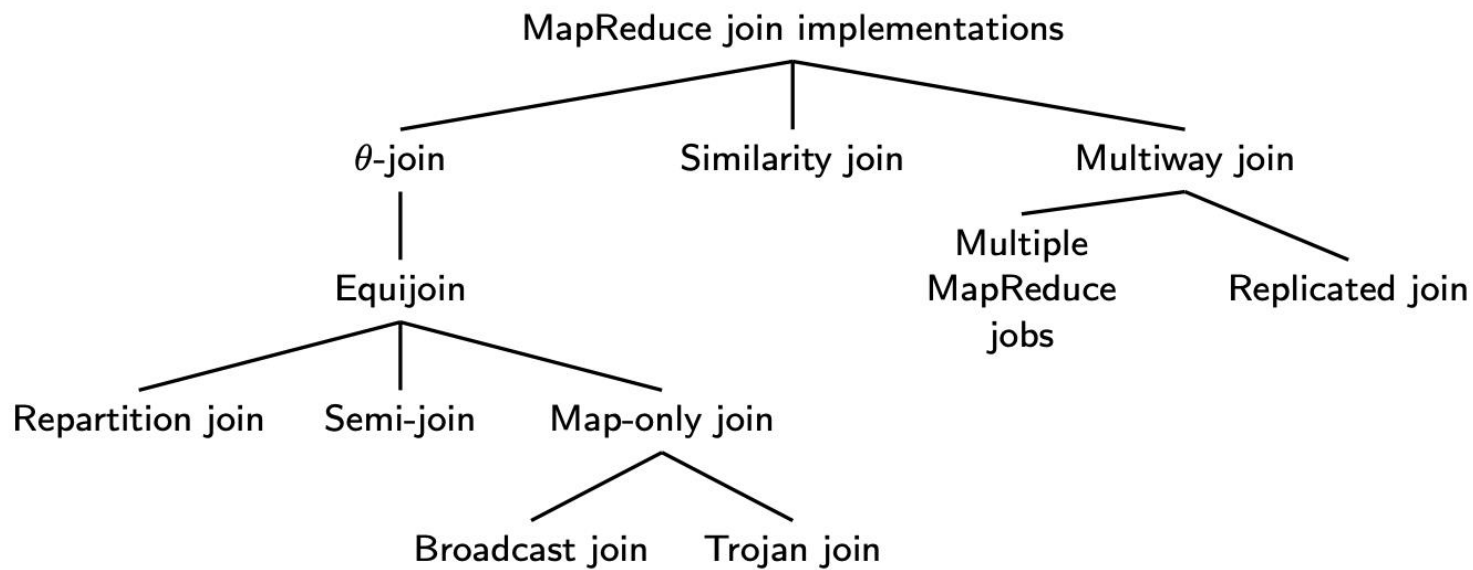
High-Level MapReduce Languages



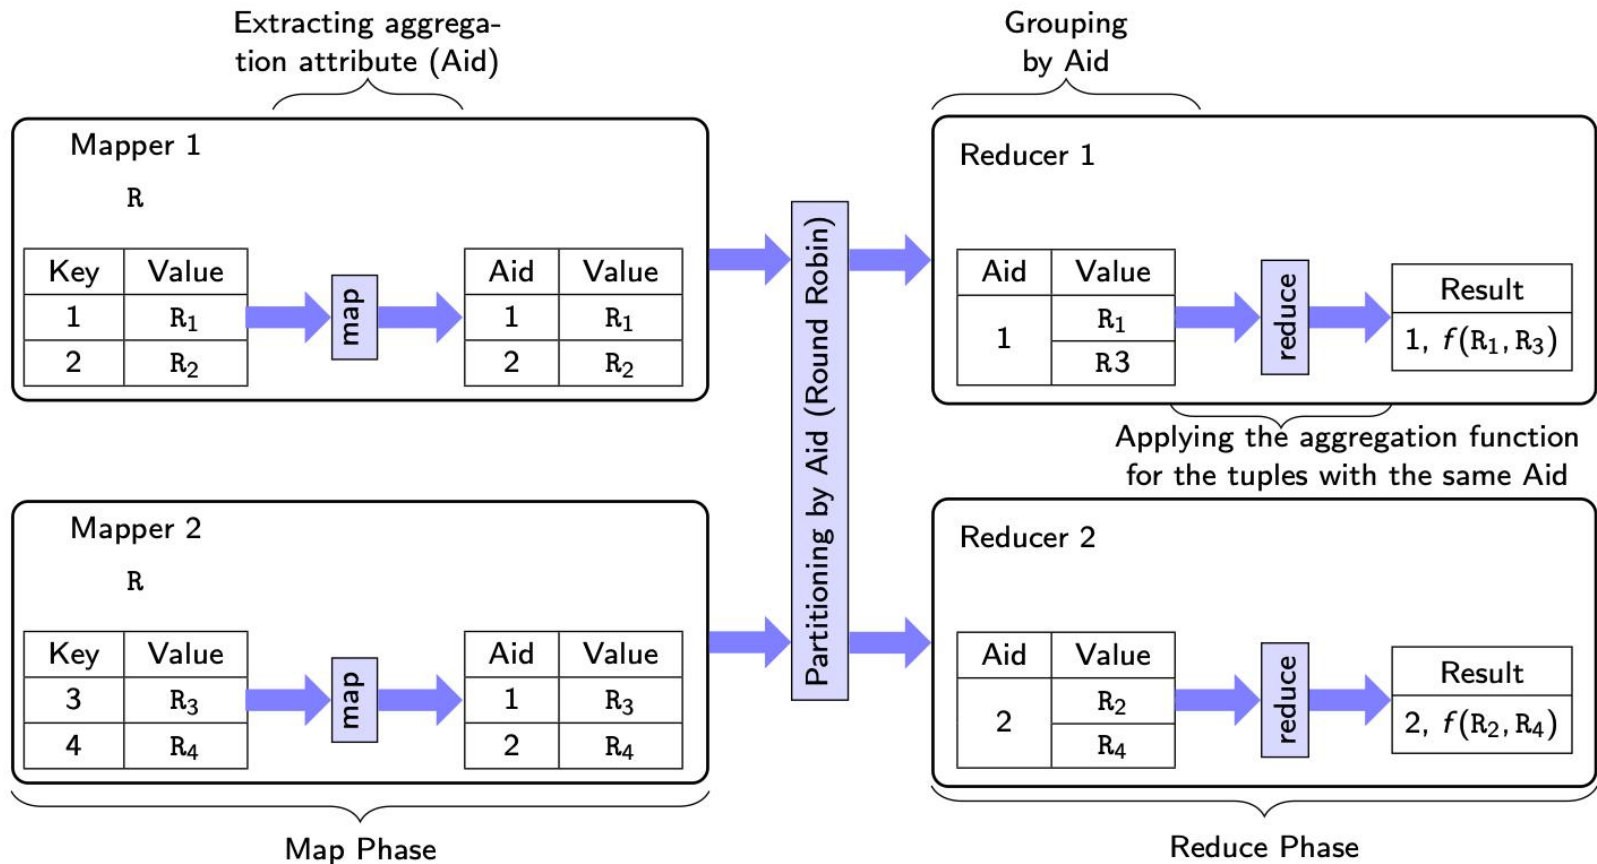
- Declarative
 - ❑ HiveQL
 - ❑ Tenzing
 - ❑ JAQL
- Data flow
 - ❑ Pig Latin
- Procedural
 - ❑ Sawzall
- Java Library
 - ❑ FlumeJava

MapReduce Implementations of DB Ops

- Select and Project can be easily implemented in the map function
- Aggregation is not difficult (see next slide)
- Join requires more work



Aggregation



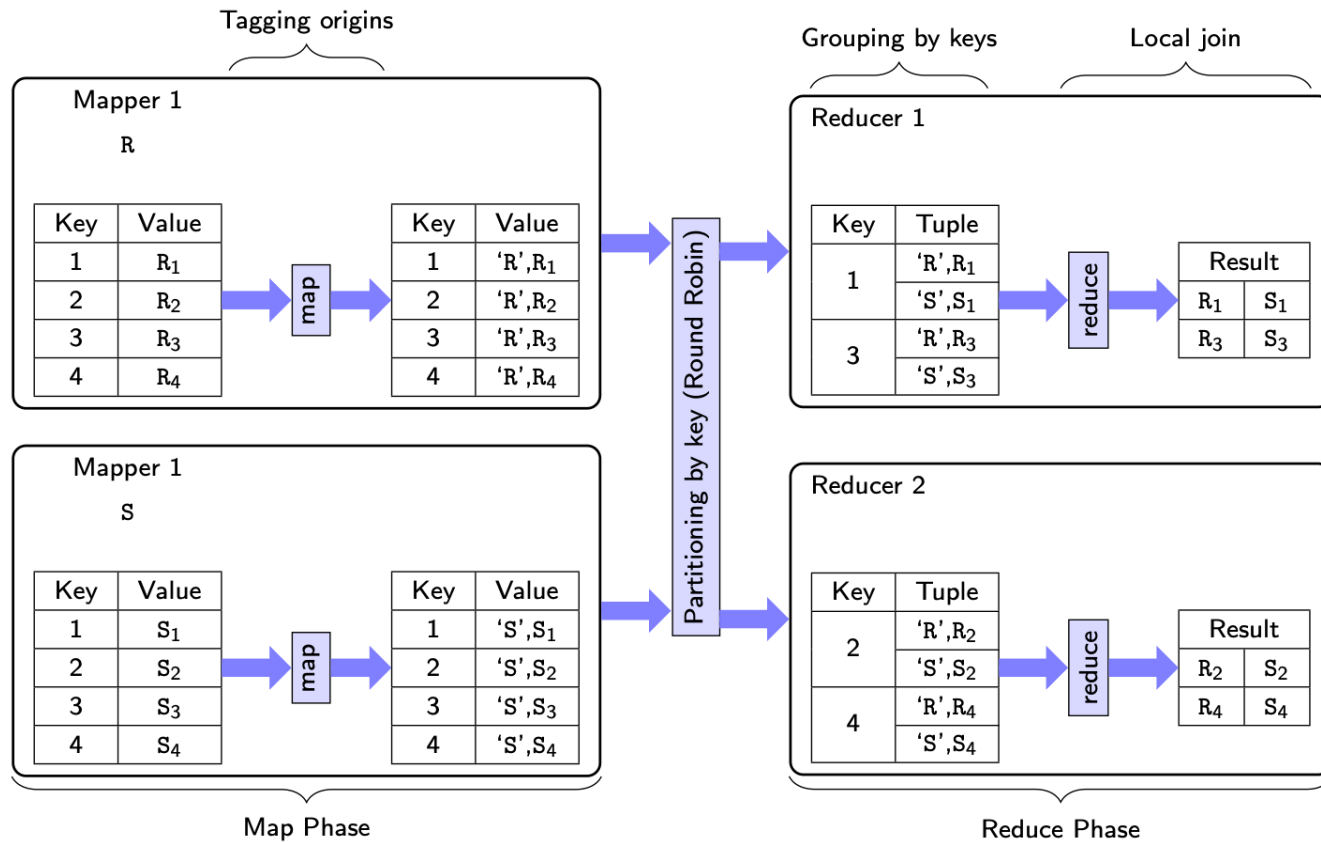
θ -Join

Baseline implementation of $R(A,B) \bowtie S(B,C)$

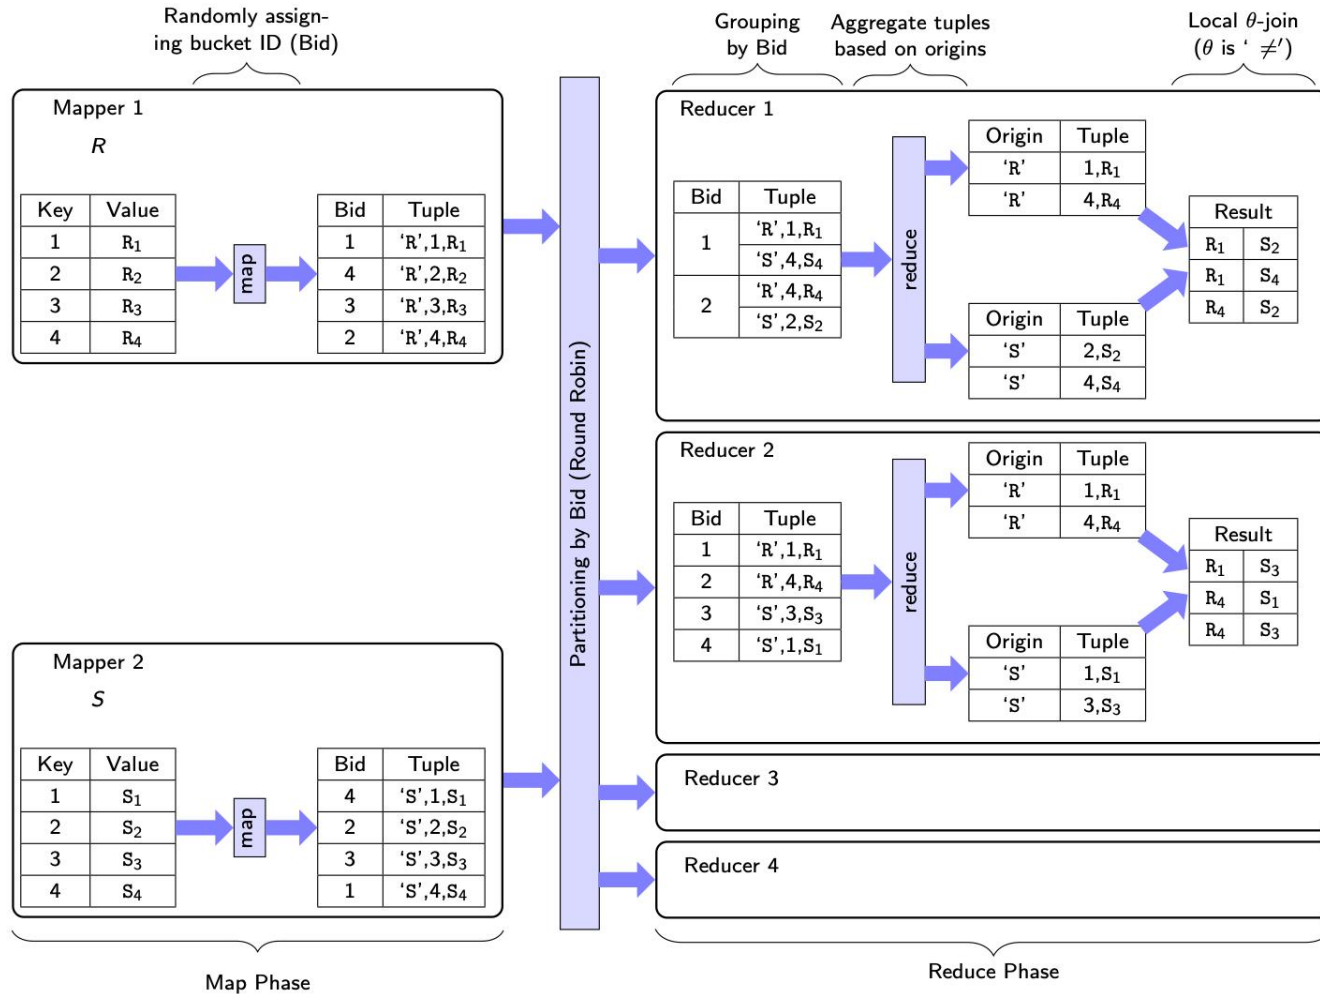
- 1) Partition R and assign each partition to mappers
- 2) Each mapper takes $\langle a,b \rangle$ tuples and converts them to a list of key-value pairs of the form $(b, \langle a,R \rangle)$
- 3) Each reducer pulls the pairs with the same key
- 4) Each reducer joins tuples of R with tuples of S

θ -Join (θ is =)

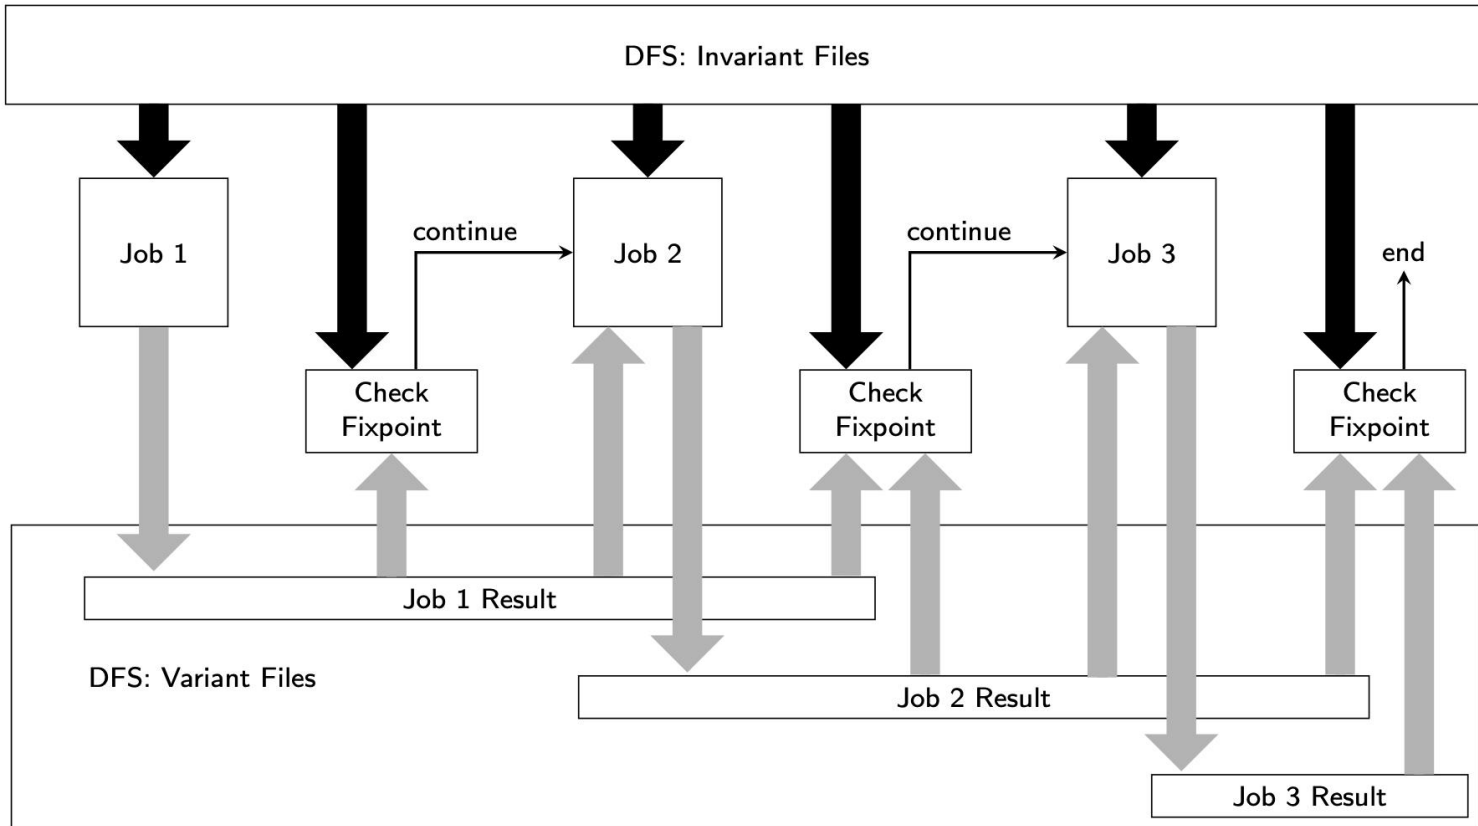
■ Repartition join



θ -Join (θ is \neq)



MapReduce Iterative Computation



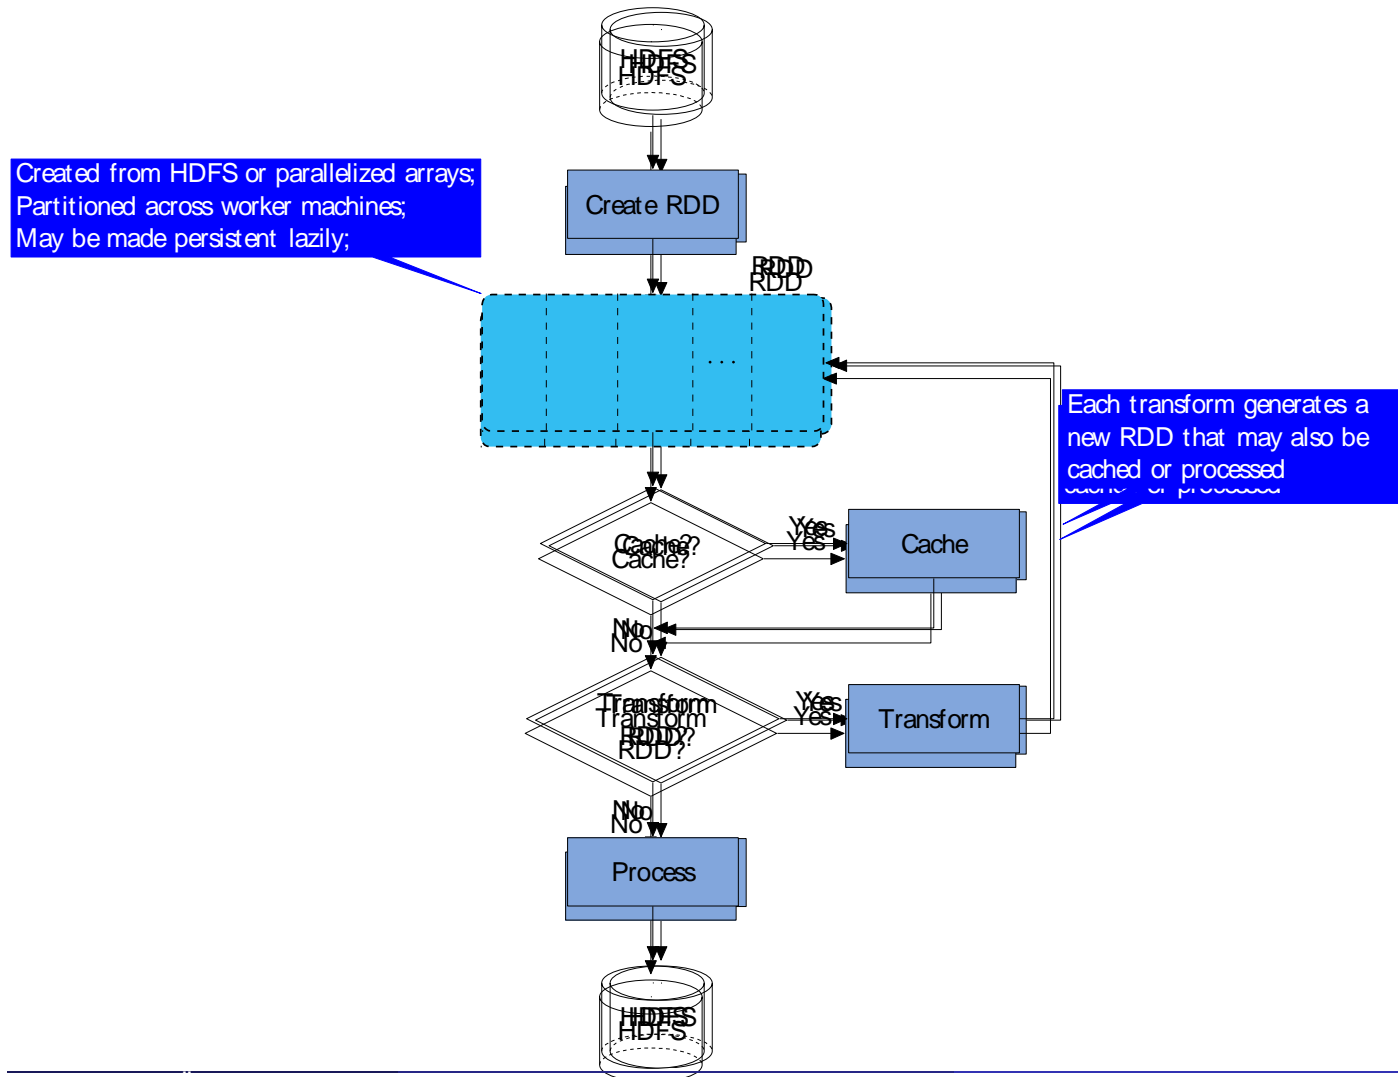
Problems with Iteration

- MapReduce workflow model is acyclic
 - ▣ Iteration: Intermediate results have to be written to HDFS after each iteration and read again
- At each iteration, no guarantee that the same job is assigned to the same compute node
 - ▣ Invariant files cannot be locally cached
- Check for fixpoint
 - ▣ At the end of each iteration, another job is needed

Spark

- Addresses MapReduce shortcomings
- Data sharing abstraction: Resilient Distributed Dataset (RDD)
 - 1) Cache working set (i.e. RDDs) so no writing-to/reading-from HDFS
 - 2) Assign partitions to the same machine across iterations
 - 3) Maintain lineage for fault-tolerance



Spark Program Flow



Outline

- Big Data Processing
 - ❑ Distributed storage systems
 - ❑ Processing platforms
 - ❑ Stream data management
 - ❑ Graph analytics

Traditional DBMS vs Streaming



■ Other differences

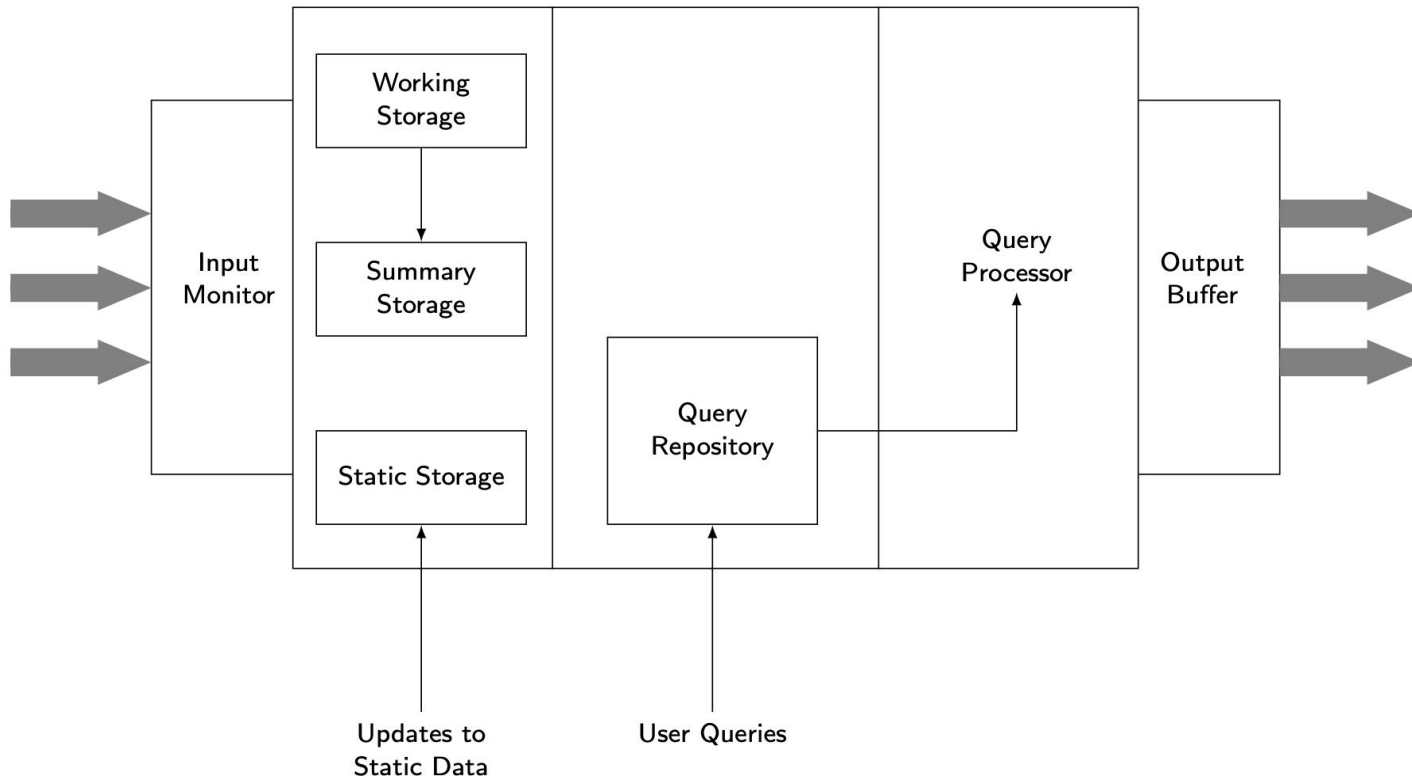
- ❑ Push-based (data-driven)
- ❑ Persistent queries

- ❑ Unbounded stream
- ❑ System conditions may not be stable

History

- **Data Stream Management System (DSMS)**
 - ❑ Typical DBMS functionality, primarily query language
 - ❑ Earlier systems: STREAM, Gigascope, TelegraphCQ, Aurora, Borealis
 - ❑ Mostly single machine (except Borealis)
- **Data Stream Processing System (DSPS)**
 - ❑ Do not embody DBMS functionality
 - ❑ Later systems: Apache Storm, Heron, Spark Streaming, Flink, MillWheel, TimeStream
 - ❑ Almost all are distributed/parallel systems
- Use **Data Stream System (DSS)** when the distinction is not important

DSMS Architecture



Stream Data Model

- Standard def: An append-only sequence of timestamped items that arrive in some order
- Relaxations
 - Revision tuples
 - Sequence of events that are reported continually (publish/subscribe systems)
 - Sequence of sets of elements (bursty arrivals)
- Typical arrival:

$\langle \text{timestamp}, \text{payload} \rangle$

 - Payload changes based on system
 - Relational: tuple
 - Graph: edge
 - ...

Processing Models

■ Continuous

- ❑ Each new arrival is processed as soon as it arrives in the system.
- ❑ Examples: Apache Storm, Heron

■ Windowed

- ❑ Arrivals are batched in windows and executed as a batch.
- ❑ For user, recently arrived data may be more interesting and useful.
- ❑ Examples: Aurora, STREAM, Spark Streaming

Window Definition

- According to the direction of endpoint movement
 - **Fixed window**: both endpoints are fixed
 - **Sliding window**: both endpoints can slide (backward or forward)
 - **Landmark window**: one endpoint fixed, the other sliding
- According to definition of window size
 - **Logical window (time-based)** – window length measured in time units
 - **Physical window (count-based)** – window length measured in number of data items
 - **Partitioned window**: split a window into multiple count-based windows
 - **Predicate window**: arbitrary predicate defines the contents of the window

Stream Query Models

- Queries are typically persistent
- They may be **monotonic** or **non-monotonic**
- Monotonic: result set always grows
 - Results can be updated incrementally
 - Answer is continuous, append-only stream of results
 - Results may be removed from the answer only by explicit deletions (if allowed)
- Non-monotonic: some answers in the result set become invalid with new arrivals
 - Recomputation may be necessary

Stream Query Languages

■ Declarative

- ❑ SQL-like syntax, stream-specific semantics
- ❑ Examples: CQL, GSQL, StreaQuel

■ Procedural

- ❑ Construct queries by defining an acyclic graph of operators
- ❑ Example: Aurora


■ Windowed languages

- ❑ `size`: window length
- ❑ `slide`: how frequently the window moves
- ❑ E.g.: `size=10min, slide=5sec`

■ Monotonic vs non-monotonic


Streaming Operators

- Stateless operators are no problem: e.g., selection, projection
- Stateful operators (e.g., nested loop join) are **blocking**
 - You need to see the entire inner operand
- For some blocking operators, non-blocking versions exist (symmetric hash join)
- Otherwise: windowed execution



Query Processing over Streams

- Similar to relational, except
 - ▣ **persistent queries**: registered to the system and continuously running
 - ▣ data pushed through the query plan, not pulled
- Stream query plan



Query Processing Issues

- Continuous execution
 - ❑ Each new arrival is processed as soon as the system gets it
 - ❑ E.g. Apache Storm, Heron
- Windowed execution
 - ❑ Arrivals are batched and processed as a batch
 - ❑ E.g. Aurora, STREAM, Spark Streaming
- More opportunities for multi-query optimization
 - ❑ E.g. Easier to determine shared subplans

Windowed Query Execution

- Two events need to be managed
 - Arrivals
 - Expirations
- System actions depend on operators
 - E.g. Join generates new result, negation removes previous result
- Window movement also affects results
 - As window moves, some items in the window move out
 - What to do to results
 - If monotonic, nothing; if non-monotonic, two options
 - Direct approach
 - Negative tuple approach

Load Management


- Stream arrival rate > processing capability
- Load shedding
 - Random
 - Semantic
- Early drop
 - All of the downstream operators will benefit
 - Accuracy may be negatively affected
- Late drop
 - May not reduce the system load much
 - Allows the shared subplans to be evaluated

Out-of-Order Processing

- Assumption: arrivals are in timestamp order
- May not hold
 - Arrival order may not match generation order
 - Late arrivals → no more or just late?
 - Multiple sources
- Approaches
 - Built-in slack
 - Punctuations

Multiquery Optimization

- More opportunity since the **persistent** queries are known beforehand
 - ▣ Aggregate queries over different window lengths or with different slide intervals
 - ▣ State and computation may be shared (usual)

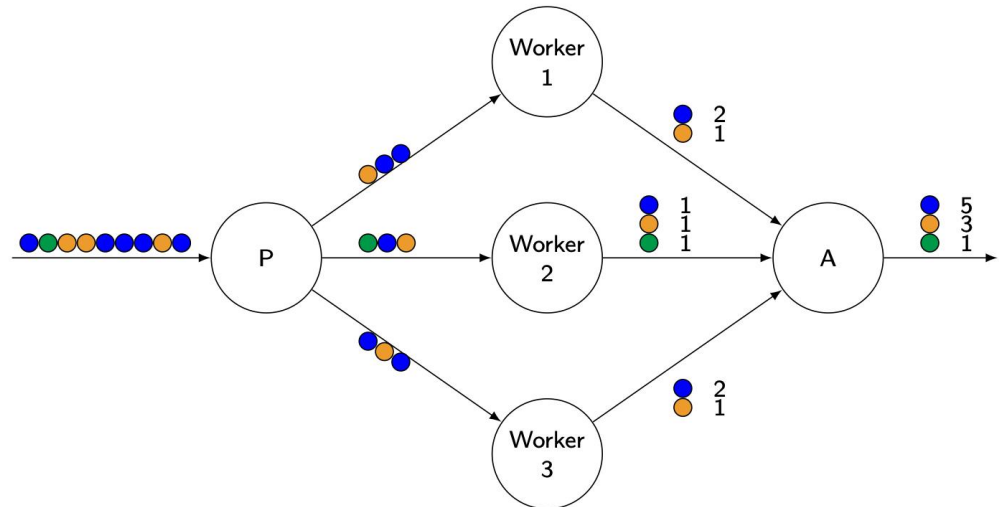


Parallel Data Stream Processing


- 1) Partitioning the incoming stream
- 2) Execution of the operation on the partition
- 3) (Optionally) aggregation of the results from multiple machines

Stream Partitioning

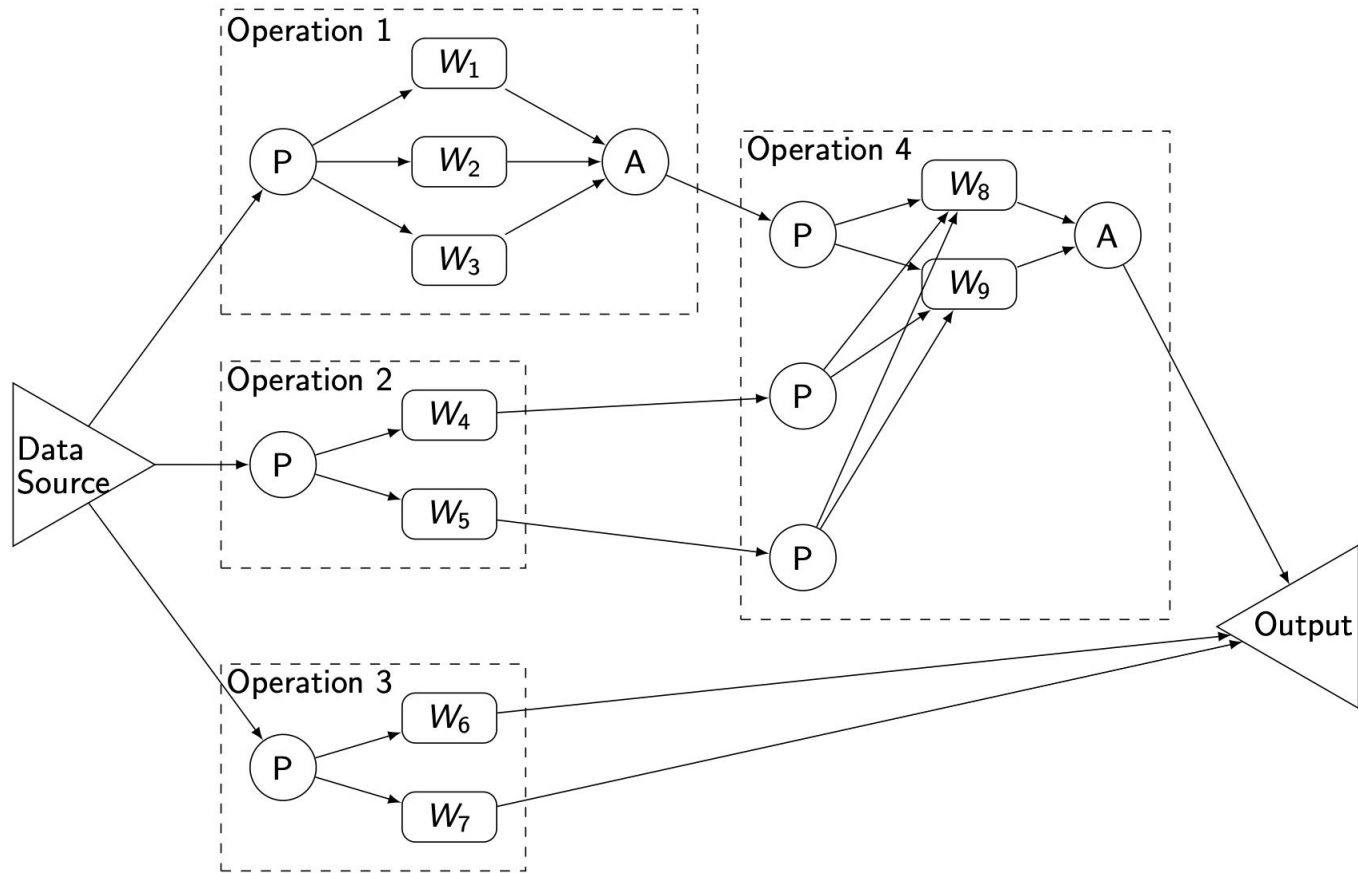
- Shuffle (round-robin) partitioning



- Hash partitioning



Parallel Stream Query Plan



Outline

- Big Data Processing
 - ▣ Distributed storage systems
 - ▣ Processing platforms
 - ▣ Stream data management
 - ▣ Graph analytics

Property Graph

- Graph $G=(V, E, D_V, D_E)$ where V is set of **vertices**, E is set of **edges**, D_V is set of **vertex properties**, D_E is set of **edge properties**
- Vertices represent entities, edges relationships among them.
- **Multigraph**: multiple edges between a pair of vertices
- **Weighted graph**: edges have weights
- **Directed** vs **undirected**

Graph Workloads

Analytical

- Multiple iterations
- Process each vertex at each iteration
- Examples
 - ❑ PageRank
 - ❑ Clustering
 - ❑ Connected components
 - ❑ Machine learning tasks

Online

- No iteration
- Usually access portion of the graph
- Examples
 - ❑ Reachability
 - ❑ Single-source shortest path
 - ❑ Subgraph matching

PageRank as Analytical Example

A web page is important if it is pointed at by other important web pages.

$$PR(P_i) = (1 - d) + d \sum_{P_j \in B_{P_i}} \frac{PR(P_j)}{|F_{P_j}|}$$


B_{P_i} : in-neighbors of P_i

F_{P_i} : out-neighbors of P_i

(let $d = 0.85$)

$$PR(P_2) = 0.15 + 0.85\left(\frac{PR(P_1)}{2} + \frac{PR(P_3)}{3}\right)$$

Recursive!...



Graph Partitioning

- Graph partitioning is more difficult than relational partitioning because of edges
- Two approaches
 - Edge-cut or vertex-disjoint
 - Each vertex assigned to one partition, edges may be replicated
 - Vertex-cut or edge-disjoint
 - Each edge is assigned to one partition, vertices may be replicated
- Objectives
 - Allocate each vertex/edge to partitions such that partitions are mutually exclusive
 - Partitions are balanced
 - Minimize edge-/vertex-cuts to minimize communication

Graph Partitioning Formalization

minimize $C(P)$

← Total communication cost
due to partitioning

subject to:

$$w(P_i) \leq \beta * \frac{\sum_{j=1}^k w(P_j)}{k}, \forall i \in \{1, \dots, k\}$$

↗ Abstract cost of
processing partition P_i

↖ Slackness for unbalanced
partitioning

- $C(P)$ and $w(P_i)$ differ for different partitionings

Vertex-Disjoint (Edge-Cut)

- Objective is to minimize the number of edge cuts
- Objective function

$$C(P) = \frac{\sum_{i=1}^k |e(P_i, V \setminus P_i)|}{|E|} \quad \text{where } |e(P_i, P_j)| = \text{\#edges between } P_i \text{ and } P_j$$

- $w(P_i)$ defined in terms of the number of vertices-per-partition


METIS Vertex-Disjoint Partitioning

- METIS is a family of algorithms
- Usually the gold standard for comparison

Given an initial graph $G_0 = (V, E)$:

- 1) Produce a hierarchy of successively coarsened graphs G_1, \dots, G_n such that $|V(G_i)| > |V(G_j)|$ for $i < j$
- 2) Partition G_n using some partitioning algorithm
 - Small enough that it won't matter what algorithm is used
- 3) Iteratively coarsen G_n to G_0 , and at each step
 - a) Project the partitioning solution on graph G_j to graph G_{j-1}
 - b) Improve the partitioning of G_0

Vertex-Disjoint Partitioning Example



Edge-Disjoint Partitioning

- Vertex-disjoint perform
 - well for graphs with low-degree vertices
 - poorly on power-law graphs causing many edge-cuts
- Edge-disjoint (vertex-cut) better for these
 - Put each edge in one partition
 - Vertices may need to be replicated – minimize these

- Objective function

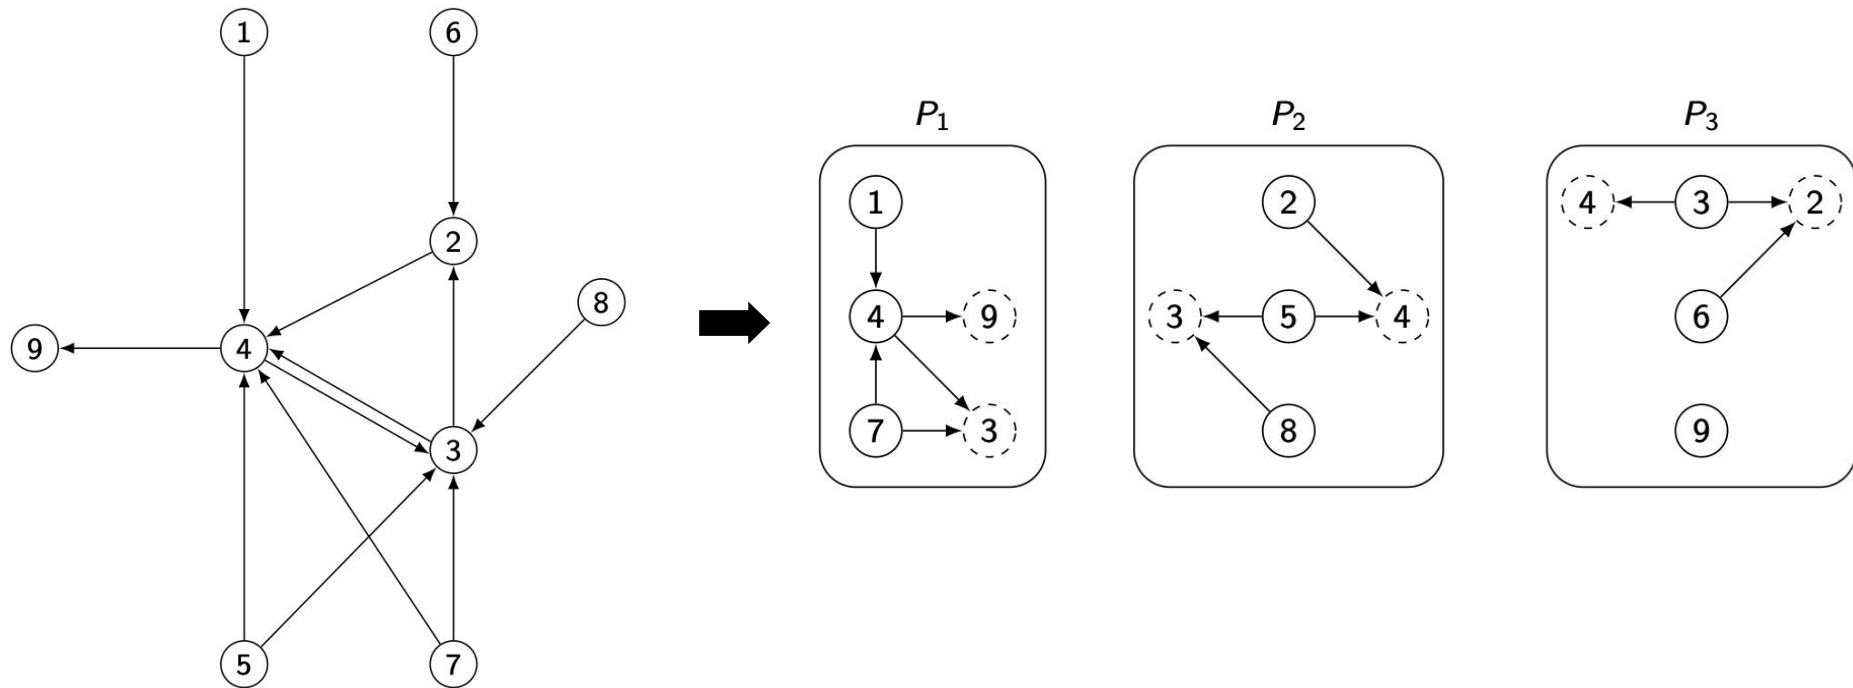
$$C(P) = \frac{\sum_{v \in V} |A(v)|}{|V|} \quad \text{where } A(v) \subseteq \{P_1, \dots, P_k\} \text{ is set of partitions in which } v \text{ exists}$$

- $w(P_i)$ is the number of edges in partition P_i

Edge-Disjoint Alternatives

- Hashing (on the ids of the two vertices incident on edge)
 - Fast and highly parallelizable
 - Gives good balance
 - But may lead to high vertex replication
- Heuristics cognizant of graph characteristics
 - Greedy: decide on allocation edge $i+1$ based on the allocation of the previous i edges to minimize vertex replication


Edge-Disjoint Partitioning Example



Can MapReduce be used for Graph Analytics?

- map and reduce functions can be written for graph analytics
 - ❑ There are works that have done this
- Graph analytics tasks are iterative
 - ❑ Recall: MapReduce is not good for iterative tasks
- Spark improves MapReduce (e.g., Hadoop) for iterative tasks
 - ❑ GraphX on top of Spark
 - ❑ Edge-disjoint partitioning
 - ❑ Vertex table & edge table as RDDs on each worker
 - ❑ Join vertex & edge tables
 - ❑ Perform an aggregation


Special-Purpose Graph Analytics Systems



???: Systems do not exist


Vertex-Centric Model

- Computation on a vertex is the focus
- “Think like a vertex”
- Vertex computation depends on its own state + states of its neighbors
- `Compute(vertex v)`
- `GetValue()`, `writeValue()`




Partition-centric (Block-centric) Model

- Computation on an entire partition is specified
- “Think like a block” or “Think like a graph”
- Aim is to reduce the communication cost among vertices

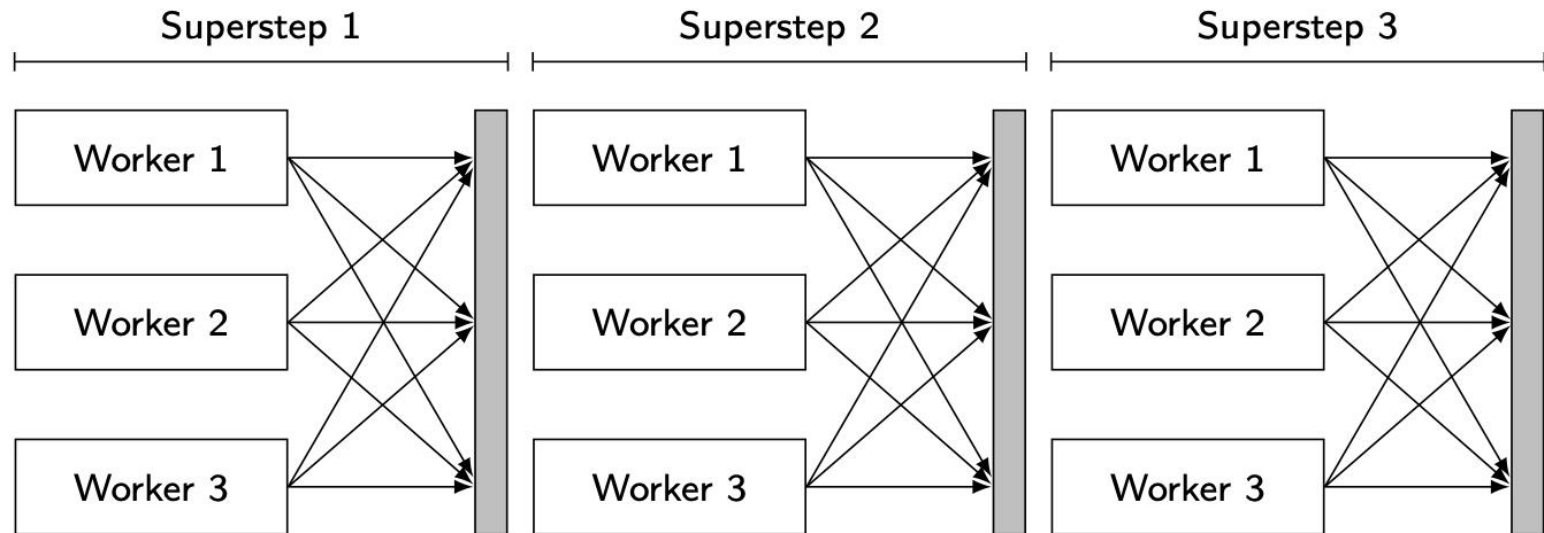


Edge-centric Model

- Computation is specified on each edge rather than on each vertex or bloc
- `Compute(edge e)`



Bulk Synchronous Parallel (BSP) Model




Each machine performs computation on its graph partition

At the end of each superstep results are **pushed** to other workers

Asynchronous Parallel (AP) Model

- Supersteps, but no communication barriers
- Uses the most recent values
- Computation in step k may be based on neighbor states of step $k-1$ (of received late) or of state k
- Consistency issues → requires distributed locking



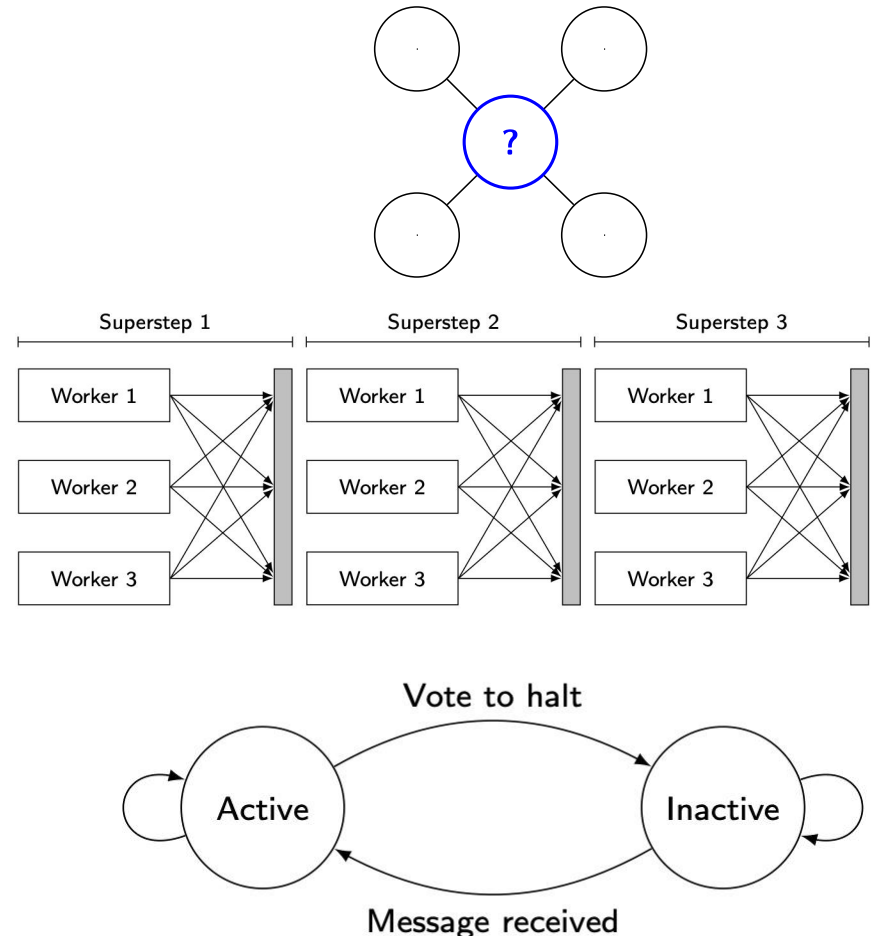
Consider vertex-centric

Gather-Apply-Scatter (GAS) Model

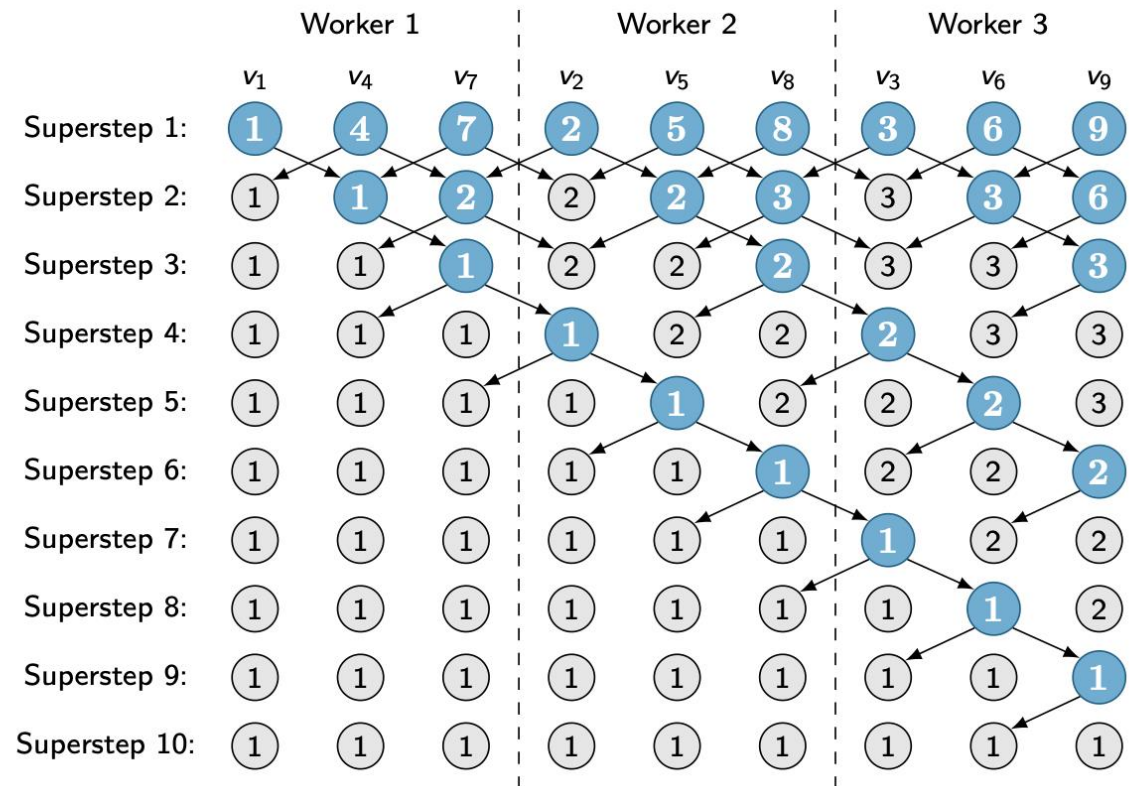

- Similar to BSP, but pull-based
- Gather: pull state
- Apply: Compute function
- Scatter: Update state
- Updates of states separated from scheduling

Vertex-Centric BSP Systems

- “Think like a vertex”
- $\text{Compute}(\text{vertex } v)$
- BSP Computation – push state to neighbor vertices at the end of each superstep
- Continue until all vertices are inactive
- Vertex state machine






Connected Components: Vertex-Centric BSP

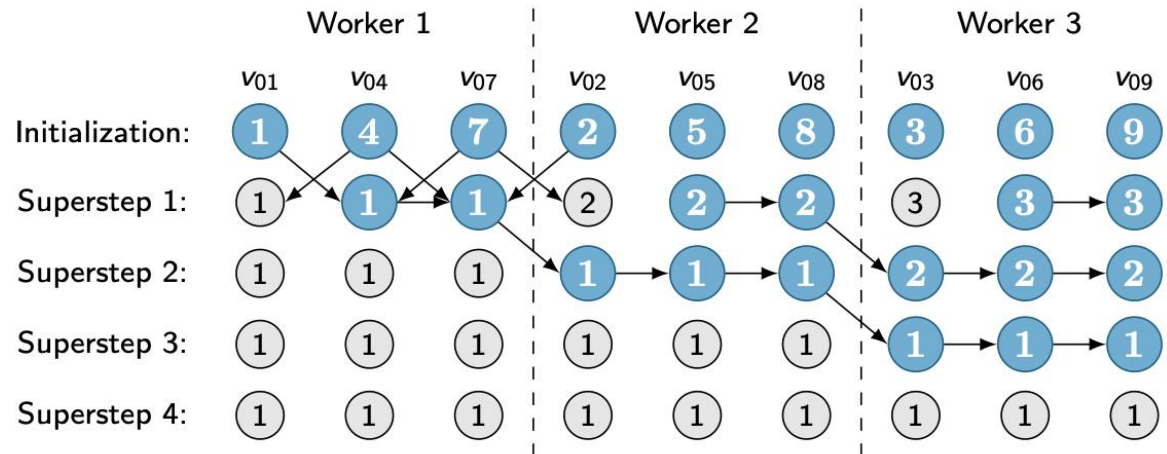



Vertex-Centric AP Systems

- “Think like a vertex”
- `Compute(vertex v)`
- Supersteps exist along with synchronization barriers, but ...
- `Compute(vertex v)` can see msgs sent in the same superstep or previous one
- Consistency of vertex states: distributed locking
- Consistency issues: no guarantee about input to `Compute()`

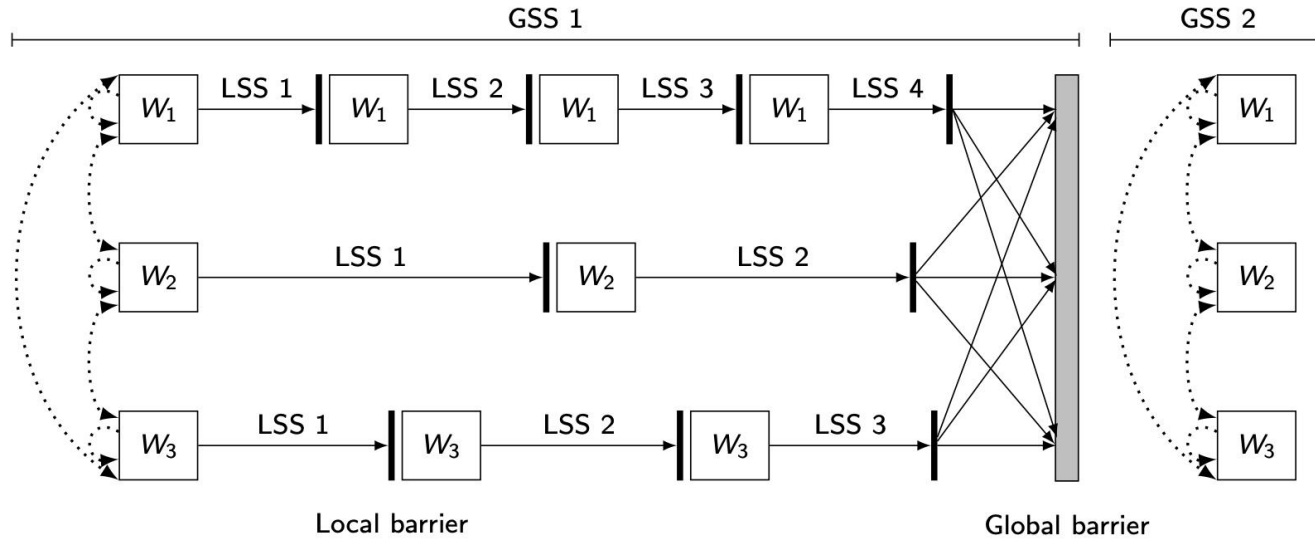


Connected Components: Vertex-Centric AP




Barrierless Asynchronous Parallel (BAP)

- Divides each global superstep into logical supersteps separated by very lightweight local barriers
- Compute() can be executed multiple times in each superstep (once-per-logical superstep)
- Synchronizes at global barriers as in AP

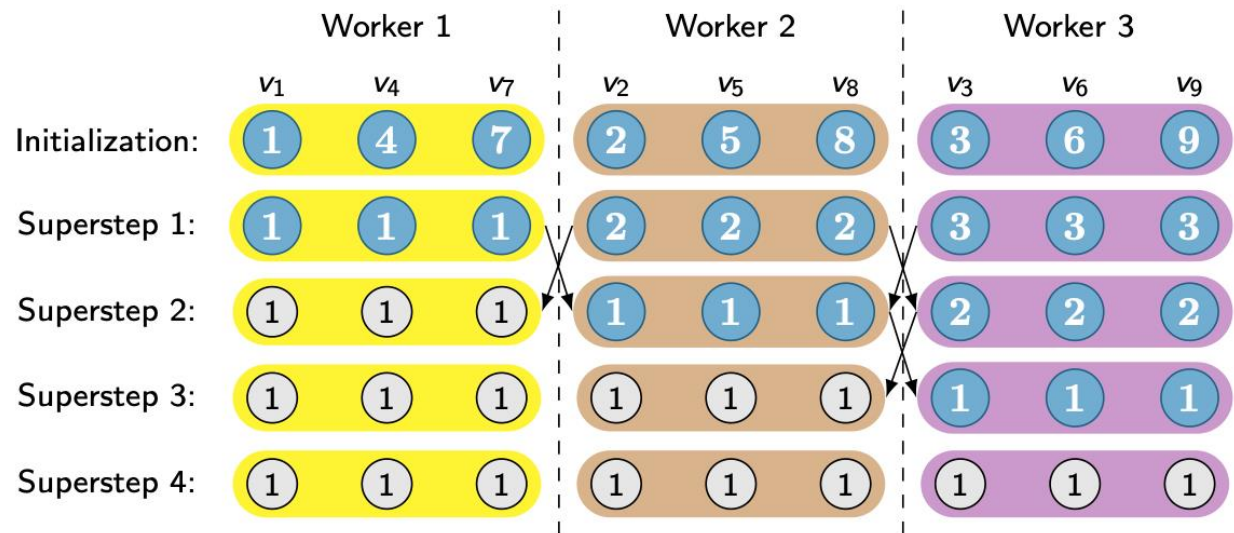



Partition- (Block-)Centric BSP Systems

- “Think like a block”; also “think like a graph”
- Reduces communication
- Exploit the partitioning of the graph
 - ▣ Message exchanges only among blocks; BSP in this case
 - ▣ Within a block, run a serial in-memory algorithm

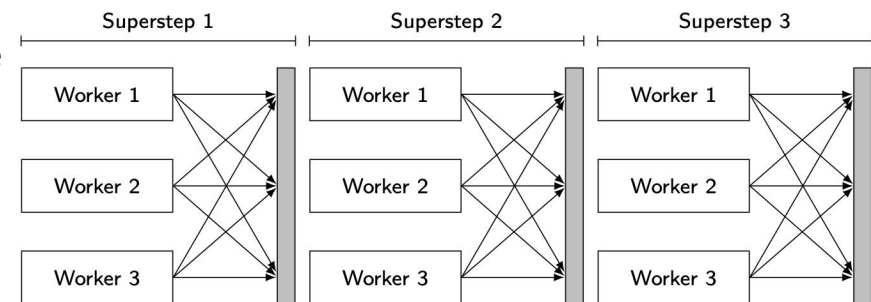



Connected Components: Block-Centric BSP



Edge-Centric BSP Systems

- Dual of vertex-centric BSP
- “Think like an edge”
- `Compute(edge e)`
- BSP Computation – push state to neighbor vertices at the end of each superstep
- Continue until all vertices are inactive
- Number of edges \gg number of vertices
 - ❑ Fewer msgs, more computation
 - ❑ No random reads



Data Lake

- Collection of raw data in native format
 - ▣ Each element has a unique identifier and metadata
 - ▣ For each business question, you can find the relevant data set to analyze it
- Originally based on Hadoop
 - ▣ Enterprise Hadoop

Advantages of a Data Lake

- Schema on read
 - ❑ Write the data as they are, read them according to a diagram (e.g. code of the Map function)
 - ❑ More flexibility, multiple views of the same data
- Multi-workload data processing
 - ❑ Different types of processing on the same data
 - ❑ Interactive, batch, real time
- Cost-effective data architecture
 - ❑ Excellent cost/performance and ROI ratio with SN cluster and open source technologies

Principles of a Data Lake

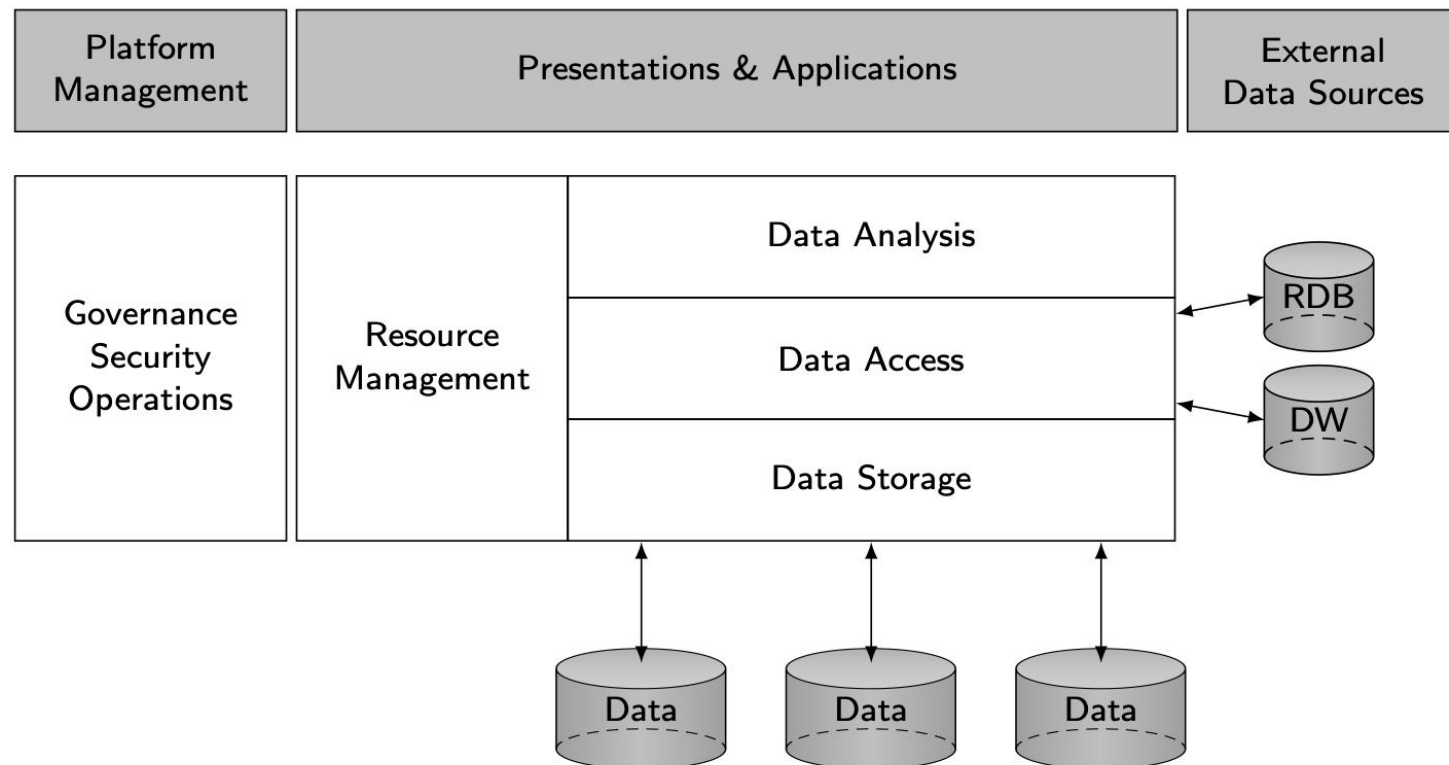
- Collect all useful data
 - Raw data, transformed data
- Dive from anywhere
 - Users from different business units can explore and enrich the data
- Flexible access
 - Different access paths to shared infrastructure
 - Batch, interactive (OLAP and BI), real-time, search,.....

Main Functions

- Data management, to store and process large amounts of data
- Data access: interactive, batch, real time, streaming
- Governance: load data easily, and manage it according to a policy implemented by the *data steward*
- Security: authentication, access control, data protection
- Platform management: provision, monitoring and scheduling of tasks (in a cluster)

Data Lake Architecture

A collection of multi-modal data stored in their raw formats



Data Lake vs Data Warehouse

Data Lake

- Shorter development process
- Schema-on-read
- Multiworkload processing
- Cost-effective architecture

Data Warehouse

- Long development process
- Schema-on-write
- OLAP workloads
- Complex development with ETL